
BOOTSTRAP CONFIDENCE INTERVALS WITH DIFFERENTIAL PRIVACY

A PREPRINT

Marcel Neunhoffer
Department of Computer Science
Boston University
Boston, MA
marceln@bu.edu

Adam Smith
Department of Computer Science
Boston University
Boston, MA
ads22@bu.edu

January 26, 2022

Abstract

Enter the text of your abstract here.

Keywords Bootstrap · Differential Privacy · Statistical Inference

1 Introduction

Here goes an introduction text

2 Our proposed method

Instead of focusing on a single quantity of interest we produce synthetic data in the form of a private (discretized) cumulative distribution function (cdf). The private cdf is then used to draw resamples and apply the same mechanism on each of the resamples. This step does not cost any additional privacy budget as it is post-processing the private cdf. We then use the resulting distribution of cdfs to calculate confidence intervals for several statistics of interest (such as quantiles, e.g. the median, the mean, the probability of observing a point in a particular region) and a confidence interval for the entire cdf.

3 Related Work

4 Experiments

We evaluate the performance of our proposed method in a set of experiments. First, we are interested in the validity of the generated confidence intervals. Second, we evaluate how different settings of the parameters in the algorithm (ϵ , B , bounds, granularity) affect the resulting confidence intervals. The results show that...

Datasets. To understand the behavior of our method under different settings we look at several data sets. First, we look at data sets where we have full control over the data generating process:

- data drawn from a standard normal distribution $\mathcal{N}(0, 1)$.
- data drawn from a mixture of two normal distributions with equal weight $f(x) = \sum_{k=1}^K \lambda_k f_k(x)$, with $\lambda = (0.5, 0.5)$, $f_1(x) = \mathcal{N}(-2, 0.25)$ and $f_2(x) = \mathcal{N}(2, 0.25)$.
- (different weights, different means, $\frac{1}{3}, \frac{2}{3}$)
- lognormal distribution




Figure 1: Summary of the data sets

Furthermore, we show results on real data and use the adult data set from the UCI Machine Learning Repository ([1]) that is derived from 1994 census data. In particular we look at the age variable.

Figure ?? summarizes the data sets. From each of the data sets we draw samples \hat{P} of different sizes (50, 100, 500, 1000, 5000), to understand the behavior of our method at different sample sizes.

(Further potential real data sets: Census (ACS and/or decennial)) PUMS data from IPUMS, Census business data)

Parameters. In our algorithm we have several parameters that influence the results. Most importantly we vary the privacy budget ϵ (0.1, 1, 10, Inf) and the number of bootstrap iterations B (1000, increase for a couple of settings). Furthermore the lower and upper bound of the cdf algorithm as well as the granularity parameter play an important role and poorly set may bias results. For the known data generating processes (normal, mixture) we set the lower bound to -4 and the upper bound to 4 and the granularity parameter to 0.01 . For the adult data we set the lower bound to 18 and the upper bound to 99 with a granularity of 0.1 . (Other sets, how to choose them?)

Evaluation. To understand the performance of our method we repeatedly (100 times) apply the method to fresh samples \hat{P} from the population data P . For each sample \hat{P} we produce confidence intervals and calculate the proportion of confidence intervals that cover the true population value (empirical coverage). Furthermore, we are interested in the length of the confidence intervals (shorter intervals with coverage close to the nominal coverage are better).

In the case of samples from the known data generating process we can compare the resulting confidence intervals to the true population values. For the adult data set (and other real data sets) we consider the full data set as our ground truth and take simple random samples from it to evaluate our proposed method. We consider the statistic of interest calculated on the entire data set as our ground truth.

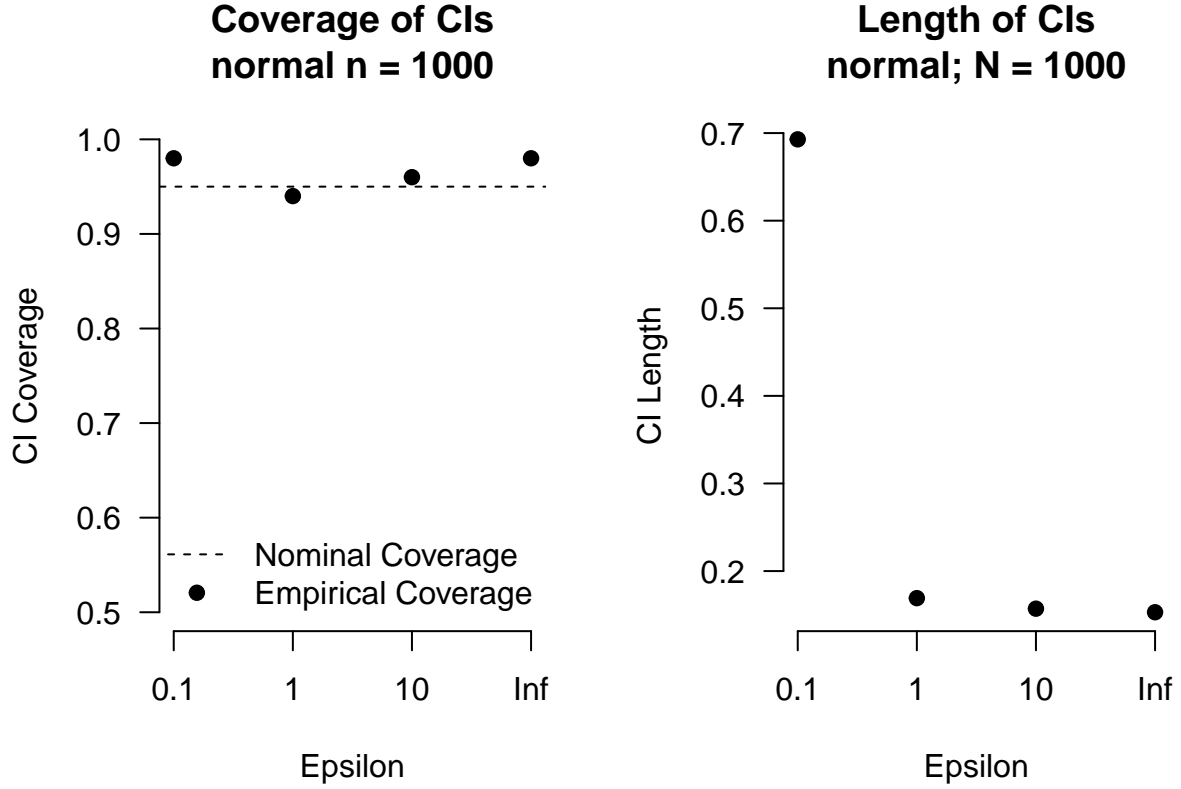


Figure 2: Results for the normal distribution

Statistics of interest. Median, mean, some percentile $\frac{1}{3}$, entire cdf \rightarrow CI around it!

5 Results

(Ratio of length to non-private one, IQR of actual distribution)

(Results to record. Bootstrap distributions, Profiling of code, Parallelization of code?)

(Combinations of parameters, other methods to get CIs, explain that pivot doesn't work in some cases and show it. True value from discretized data.)

6 Notes

Different possibilities to get quantiles. So far we used the percentile method.

The range and granularity need to be set properly (how?), otherwise truncation and (too coarse) discretization can bias results.




Figure 3: Results for the mixture distribution

7 Instructions

You can use directly LaTeX command or Markdown text.

LaTeX command can be used to reference other section. See Section ???. However, you can also use **bookdown** extensions mechanism for this.

7.1 Headings: second level

You can use equation in blocks

$$\xi_{ij}(t) = P(x_t = i, x_{t+1} = j | y, v, w; \theta) = \frac{\alpha_i(t) a_{ij}^{w_t} \beta_j(t+1) b_j^{v_{t+1}}(y_{t+1})}{\sum_{i=1}^N \sum_{j=1}^N \alpha_i(t) a_{ij}^{w_t} \beta_j(t+1) b_j^{v_{t+1}}(y_{t+1})}$$

But also inline i.e $z = x + y$

7.1.1 Headings: third level

Another paragraph.

8 Examples of citations, figures, tables, references

You can insert references. Here is some text (Kour and Saabne 2014b, 2014a) and see Hadash et al. (2018).

The documentation for **natbib** may be found at

You can use custom blocks with LaTeX support from **rmarkdown** to create environment.




Figure 4: Results for the adult data

<http://mirrors.ctan.org/macros/latex/contrib/natbib/natnotes.pdf%7D>

Of note is the command `\citet`, which produces citations appropriate for use in inline text.

You can insert LaTeX environment directly too.

```
\citet{hasselmo} investigated\dots
```

produces

Hasselmo, et al. (1995) investigated...

<https://www.ctan.org/pkg/booktabs>

8.1 Figures

You can insert figure using LaTeX directly.

See Figure ???. Here is how you add footnotes. [[^]Sample of the first footnote.]

But you can also do that using R.

```
plot(mtcars$mpg)
```

You can use **bookdown** to allow references for Tables and Figures.

8.2 Tables

Below we can see how to use tables.




Figure 5: Sample figure caption.




Figure 6: Another sample figure

Table 1: Sample table title

Part		
Name	Description	Size (μm)
Dendrite	Input terminal	~ 100
Axon	Output terminal	~ 10
Soma	Cell body	up to 10^6

See awesome Table~1 which is written directly in LaTeX in source Rmd file.

You can also use R code for that.

```
knitr::kable(head(mtcars), caption = "Head of mtcars table")
```

Table 2: Head of mtcars table

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

8.3 Lists

- Item 1
- Item 2
- Item 3

Hadash, Guy, Einat Kermany, Boaz Carmeli, Ofer Lavi, George Kour, and Alon Jacovi. 2018. “Estimate and Replace: A Novel Approach to Integrating Deep Neural Networks with Existing Applications.” *arXiv Preprint arXiv:1804.09028*.

Kour, George, and Raid Saabne. 2014a. “Fast Classification of Handwritten on-Line Arabic Characters.” In *Soft Computing and Pattern Recognition (SoCPaR), 2014 6th International Conference of*, 312–18. IEEE.

———. 2014b. “Real-Time Segmentation of on-Line Handwritten Arabic Script.” In *Frontiers in Handwriting Recognition (ICFHR), 2014 14th International Conference on*, 417–22. IEEE.

References

- [1] Dheeru Dua and Casey Graff. UCI machine learning repository, 2017.